

BASIC ALGORITHMS

Algorithms for small problems

Diego Klabjan
Professor, Industrial Engineering and Management Sciences

Northwestern | McCORMICK SCHOOL OF
ENGINEERING

Outline

- Enumeration
- Value iteration
- Policy iteration
 - Value of policy
- Hybrid

ENUMERATION

Bellman's Optimality Equation

- $V_t(s_t) = \max_{a_t \in \mathcal{A}} [r(s_t, a_t) + \gamma E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)} V_{t+1}(s_{t+1}|s_t)]$
- Compute value functions recursively
 - Training (planning)
- Given computed value functions
 - 'Measure' state
 - Solve optimization problem

$$\max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$$

- Often not many actions – enumerate
- Often deterministic system – no expectation

Enumeration


- For $t=T$ down to 0
 - For each possible state s_t
 - Compute

$$V_t(s_t) = \max_{a_t \in \mathcal{A}} [r(s_t, a_t) + \gamma E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)} V_{t+1}(s_{t+1}|s_t)]$$

- Works if
 - Small number of states
 - Small number of actions
 - Somehow cope with expectation
- Three courses of dimensionality

Enumeration Example

- 5 drones in a room of size 30 x 30 x 30 (feet)
 - Discretize by 6 inches
- Guide them to avoid collisions
 - Each one has its own destination
- Number of possible states $(60 \cdot 60 \cdot 60)^5 \approx 4.7 \cdot 10^{26}$
 - Small room, imagine a warehouse
 - Not capturing angle
- Actions
 - $27^3 \approx 20,000$
 - Tractable



<http://www.m759.net/wordpress/?s=galois+cube>

VALUE ITERATION

Enumeration

- Episode length vary or is infinite
 - Enumeration does not work
 - Vast majority of practical problems
- Recursively defined optimality equation

$$V(s) = \max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$$

- Can be shown

$$V_t^T(s_t) = \max_{a_t \in \mathcal{A}} [r(s_t, a_t) + \gamma E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)} V_{t+1}^T(s_{t+1}|s_t)]$$

$$V(s) = \lim_{T \rightarrow \infty} V_0^T(s)$$

Value Iteration

- $V(s) = \max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$
- Assume right-hand side known
 - Use approximate V
 - Can compute left-hand side
 - Gives better approximation of V

Value Iteration

- For $k = 0, 1, 2, \dots$
 - For each possible state s
 - Compute


$$V_{k+1}(s) = \max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V_k(\bar{s}|s)]$$

- If discount factor less than 1 and everything is finite
 - Convergence (pointwise) to optimal value function
- Same pitfalls as enumeration
- No explicit policy

Value Function and Policy

- $\pi(s) = \underset{a \in \mathcal{A}}{\operatorname{argmax}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$
 - If V optimal, π is optimal
- $V(s) = r(s, \pi(s)) + \gamma E_{\bar{s} \sim p(\bar{s}|s, \pi(s))} V^\pi(\bar{s}|s)$
 - Must know V^π
 - If π is optimal, V is optimal

Value Iteration



$$V_{k+1}(s) = \max_{a \in \mathcal{A}} \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_k(s') \right)$$


$$R_s^a = r(s, a)$$

$$R^a = (r(s, a))_{s \in \mathcal{S}}$$

$$V_{k+1} = \max_{a \in \mathcal{A}} R^a + \gamma P^a V_k$$

Apply for each state

Example



- Actions: up, down, left, or right.
 - 0.7 chance of going one step in the desired direction
 - 0.1 chance of going one step in any of the other three directions
- Bump into the outside wall
 - Penalty of 1 (reward of -1)
 - Agent does not actually move
- There are four rewarding states
 - In each of these states, the agent gets the reward after it carries out an action in that state
 - Not when it enters the state.

https://artint.info/html/ArtInt_224.html#gridworld-ex

Example

- 9 cells around +10
- Discount 0.9
- Start with value function of 0

0	0	-0.1
0	10	-0.1
0	0	-0.1

0	6.3	-0.1
6.3	9.8	6.2
0	6.3	-0.1

4.5	6.2	4.4
6.2	9.7	6.6
4.5	6.1	4.4

Demo of Example

- <http://www.cs.ubc.ca/~poole/demos/mdp/vi.html>

Example: Shortest Path

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

V_1

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

V_2

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

V_3

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

V_4

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

V_5

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

V_6

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

V_7

Other Applications of Value Iteration

- Shortest Path can be solved by value iteration
- Other algorithms
 - Levensthein distance
 - String algorithms
 - String alignment
 - Dynamic time warping
 - Generalization of Levensthein
 - Graphical models
 - Viterbi algorithm

POLICY ITERATION

Evaluating Policy

- Given policy π find V^π

$$V^\pi(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma V^\pi(\bar{s}|s)]$$

- Can use similar idea to value iteration
- Given approximate right-hand side
 - Find better left-hand side by using the equation


Iterative Policy Evaluation

- Problem
 - Evaluate given policy π
- Solution: iterative application of Bellman expectation equation
- $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$
- For $k = 0, 1, 2, \dots$
 - For each possible state s compute

$$v_{k+1}(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma v_k(\bar{s}|s)]$$

- Convergence to V^π can be proven

Iterative Policy Evaluation



$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_k(s') \right)$$

$$\pi(a|s) = P[a = \pi(s)]$$

$$\mathbf{v}_{k+1} = \mathbf{R}^\pi + \gamma \mathbf{P}^\pi \mathbf{v}_k$$

Apply for each state

Evaluate Policy

$$v_{k+1} = R^\pi + \gamma P^\pi v_k$$

- Limit $k \rightarrow \infty$

$$v^\pi = R^\pi + \gamma P^\pi v^\pi$$

$$v^\pi = (I - \gamma P^\pi)^{-1} R^\pi$$

- Algorithm is a way to compute the inverse
 - Inverse exists if discount less than 1
- $R^\pi = \left(E_{a \sim \pi(a|s)} r(s, a) \right)_s = \left(\sum_{a \in \mathcal{A}} \pi(a|s) r(s, a) \right)_s$
- $P^\pi = \left(\sum_{a \in \mathcal{A}} \pi(a|s) P_{ss'}^a \right)_{s,s'}$

Example of Evaluating Random Policy



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$
on all transitions

- Undiscounted episodic MDP ($\gamma = 1$)
- Nonterminal states 1, ..., 14
- Shaded squares terminal nodes
- Actions leading out of the grid leave state unchanged
- Reward is -1 until the terminal state is reached
- Agent follows uniform random policy
 - Each action selected with same probability

Example of Evaluating Random Policy

$k = 0$

v_k for random policy

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

greedy policy
with respect to
 v_k

	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	

random policy

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕	↕
↑	↕	↕	↕
↕	↕	↕	↓
↕	↕	→	

24

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕
↑	↖	↕	↓
↑	↕	↗	↓
↕	→	→	

Reinforcement Learning

Example of Evaluating Random Policy

$k = 3$

$k = 10$

$k = \infty$

v_k for random policy

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

greedy policy with respect to v_k

	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	


	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	

	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	


Policy Iteration

- Given a policy π
 - Evaluate policy π
$$V^\pi(s) = \mathbb{E}[r_t + \gamma r_{t+1} + \dots | S_t = s]$$
 - Improve the policy by acting greedily with respect to V^π
$$\pi' = \text{greedy}(V^\pi)$$
- This process of policy iteration always converges to π^* - optimal policy
 - Finite cardinality assumptions

Policy Iteration



Policy Evaluation Estimate v_π
 Iterative policy evaluation
Policy Improvement Generate $\pi' \geq \pi$
 Greedy policy improvement



Policy Improvement

- Consider a deterministic policy, $a = \pi(s)$
- We can *improve* the policy by acting greedily
$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q^\pi(s, a)$$
- Improves the value from any state s over one step
$$Q^\pi(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^\pi(s, a) \geq Q^\pi(s, \pi(s)) = V^\pi(s)$$
- Improves the value function, $V^{\pi'}(s) \geq V^\pi(s)$
$$\begin{aligned} V^\pi(s) &\leq Q^\pi(s, \pi'(s)) = \mathbb{E}_{\pi'}[r_t + \gamma V^\pi(S_{t+1}) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[r_t + \gamma Q^\pi(S_{t+1}, \pi'(S_{t+1})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[r_t + \gamma r_{t+2} + \gamma^2 Q^\pi(S_{t+2}, \pi'(S_{t+2})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[r_t + \gamma r_{t+2} + \dots | S_t = s] = V^{\pi'}(s) \end{aligned}$$

Policy Improvement

- If improvements stop

$$Q^{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^{\pi}(s, a) = Q^{\pi}(s, \pi(s)) = V^{\pi}(s)$$

- Bellman optimality equation has been satisfied

$$V^{\pi}(s) = \max_{a \in \mathcal{A}} Q^{\pi}(s, a)$$

- Conclusion $V^{\pi}(s) = V^*(s)$ for all $s \in \mathcal{S}$
- π is an optimal policy


Modified Policy Iteration

- Stopping criteria
 - ϵ -convergence of value function
 - Stop after K iterations evaluating the policy
- In the example, $K = 3$ was sufficient to achieve optimal policy

Generalized Policy Iteration

- Loop
 - For $k = 0, 1, 2, \dots, K$ // K iterations to evaluate the policy
 - For each possible state s compute
$$v_{k+1}(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma v_k(\bar{s}|s)]$$
 - Improve the policy by acting greedily with respect to v_{K+1}
$$\pi' = \text{greedy}(v_{K+1})$$
- The inner loop approximately computes the inverse of the matrix in $\mathbf{v}^\pi = (\mathbf{I} - \gamma \mathbf{P}^\pi)^{-1} \mathbf{R}^\pi$
- $K=0$
 - Value iteration

Generalized Policy Iteration




Policy Evaluation Estimate V^π

Any policy evaluation algorithm

Policy Improvement Generate $\pi' \geq \pi$

Any policy improvement algorithm



- Value iteration

- Per iteration time low
- Needs more iterations

- Policy iteration

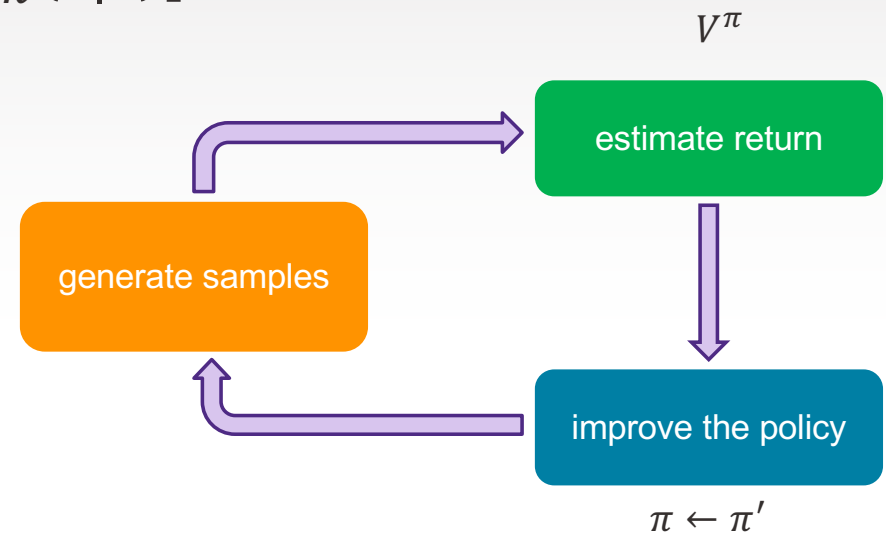
- Per iteration time high
 - Controlled by K
- Needs fewer iterations
- More flexible

Weaker convergence assumptions for policy iteration

Trade-off


Expectation

- $$v_{k+1}(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma v_k(\bar{s}|s)]$$
- Number of actions large
 - Sample actions
- For every state s
 - For $n = 1, 2, \dots, N$
 - Sample $a_n \sim \pi(a_n|s)$
 - Sample $s'_n \sim p(s'_n|s, a_n)$
 - $$v_{k+1}(s) = \frac{1}{N} \sum_{n=1}^N [r(s, a_n) + \gamma v_k(s'_n)]$$
- In practice $N=1$



Large State Space and Value Iteration

- Tabular value function
 - One big array with one entry per state
 - Not scalable
- Approximate value function
 - Neural network function $V: \mathcal{S} \rightarrow \mathbb{R}$
 - Can be any parametric function
 - No softmax
 - Output a single value



Functional Approximation of V

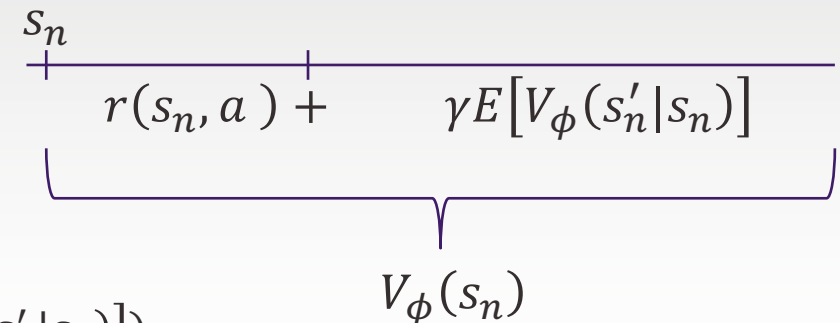
- Updating the value function approximation
- At state s
 - Have reward plus future based on approximate value function
 - Based on optimal action
 - Have value of approximate value function
- Match them
 - L2 loss

fitted value iteration algorithm:

1. set $y_i \leftarrow \max_{a_i} (r(s_i, a_i) + \gamma E[V_\phi(s'_i)])$
2. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$

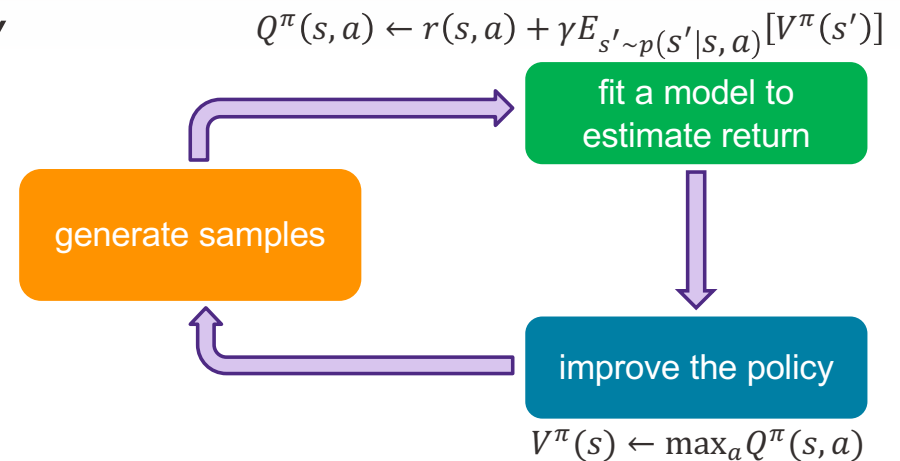
Putting it all Together

- For $k = 1, 2, \dots$
 - For $n = 0, 1, 2, \dots, N$
 - Sample s_n
 - For $n = 0, 1, 2, \dots, N$
 - Compute $y_n = \max_a (r(s_n, a) + \gamma E[V_\phi(s'_n | s_n)])$
 - Set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_n \|V_\phi(s_n) - y_n\|^2$
- Last step
 - Standard gradient optimization by back propagation
- Drawback
 - No policy improvement


$$\overbrace{r(s_n, a) + \gamma E[V_\phi(s'_n | s_n)]}^{V_\phi(s_n)}$$

Revised

- Value function as neural network
- Keep track of Q-factor at samples generated
- Compute $Q^\pi(s, a)$ at generated samples
 - Samples generated based on incumbent policy
 - Use relationship between Q and V
- Update policy in greedy manner from Q



Algorithm

- For $k = 1, 2, \dots$
 - For $n = 0, 1, 2, \dots, N$
 - Sample $s_n, a_n \sim \pi(a_n | s_n)$
 - For $n = 0, 1, 2, \dots, N$
 - $Q^\pi(s_n, a_n) \leftarrow r(s_n, a_n) + \gamma E_{s' \sim p(s' | s_n, a_n)} [V_\phi(s')]$
 - Compute $y_n = \max_a (r(s_n, a) + \gamma E_{s' \sim p(s' | s_n, a)} [V_\phi(s')])$
 - Set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_n \|V_\phi(s_n) - y_n\|^2$
 - $\pi(s) \leftarrow \operatorname{argmax}_a Q^\pi(s, a)$

Issues

- How to generate samples from state space?
 - Uniformly at random – unclear for large state space
- Episodes present
 - Sample from episodes
 - Never sample a new state
 - Yes as next state
- Why not parametric policy?
 - Update of parameters not clear
 - Needs further tricks

Unknown Transition Function

- Observe only episodes – imitation learning
- Functional approximation of transition function not possible
 - Output state

fitted value iteration algorithm:

1. set $y_i \leftarrow \max_{a_i} (r(s_i, a_i) + \gamma E[V_\phi(s'_i)])$
 2. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$
- Need to know outcomes for different actions!

policy iteration algorithm

1. evaluate $Q^\pi(s, a)$
2. set $\pi \leftarrow \pi'$

$$\pi'(a_t | s_t) = \begin{cases} 1 & \text{if } a_t = \operatorname{argmax}_{a_t} Q^\pi(s_t, a_t) \\ 0 & \text{otherwise} \end{cases}$$

policy evaluation:

$$V^\pi(s) \leftarrow r(s, \pi(s)) + \gamma E_{s' \sim p(s' | s, \pi(s))} [V^\pi(s')]$$

$$Q^\pi(s, a) \leftarrow r(s, a) + \gamma E_{s' \sim p(s' | s, \pi(s))} [Q^\pi(s', \pi(s'))]$$

Functional Approximation of Q

policy iteration:

1. evaluate $V^\pi(s)$
2. set $\pi \leftarrow \pi'$


fitted value iteration algorithm:

1. set $y_i \leftarrow \max_{a_i}(r(s_i, a_i) + \gamma E[V_\phi(s'_i)])$
2. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$

fitted Q iteration algorithm:

1. set $y_i \leftarrow r(s_i, a_i) + \gamma E[V_\phi(s'_i)]$ ← approximate $E[V(s'_i)] \approx \max_{a'} Q_\phi(s'_i, a')$
2. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

Doesn't require simulation of actions!



Value Iteration with Fitted Q-factor

- Observed data are trajectories
 - Formally, $U = \{(s_i, a_i, s'_i, r_i) | i \in N\}$
- Loop
 - Sample $S \subseteq U$
 - For $i \in S$
 - $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
 - Set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_{i \in S} \|Q_\phi(s_i, a_i) - y_i\|^2$
- Only remaining drawback how to compute max

Analysis

full fitted Q-iteration algorithm:

1. dataset $\{(s_i, a_i, s'_i, r_i)\}$
2. set $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$ ← this max improves the policy (tabular case)
3. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$ ↖ error \mathcal{E}

$$\mathcal{E} = \frac{1}{2} E_{(s,a) \sim \beta} \left[Q_\phi(s, a) - \left[r(s, a) + \gamma \max_{a'} Q_\phi(s', a') \right] \right]^2$$

- If $\mathcal{E} = 0$, then $Q_\phi(s, a) = r(s, a) + \gamma \max_{a'} Q_\phi(s', a')$
 - This is an *optimal* Q-function, corresponding to optimal policy π^* :

$$\pi^*(a_t | s_t) = \begin{cases} 1 & \text{if } a_t = \operatorname{argmax}_{a_t} Q_\phi(s_t, a_t) \\ 0 & \text{otherwise} \end{cases}$$

Online Version

full fitted Q-iteration algorithm:

1. collect dataset $\{(s_i, a_i, s'_i, r_i)\}$ using some policy
2. set $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'} Q_\phi(s'_i, a')$
3. set $\phi \leftarrow \operatorname{argmin}_\phi \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

online Q iteration algorithm:

1. take some action a_i and observe (s_i, a_i, s'_i, r_i)
2. $y_i = r(s_i, a_i) + \gamma \max_{a'} Q_\phi(s'_i, a')$
3. $\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(s_i, a_i)(Q_\phi(s_i, a_i) - y_i)$

