

BASIC ALGORITHMS

Q- and TD-learning

Diego Klabjan
Professor, Industrial Engineering and Management Sciences

Northwestern | McCORMICK SCHOOL OF
ENGINEERING

Outline

- Q-learning
 - Variants
 - Convergence
- TD-learning
 - Monte-Carlo sampling
 - TD(λ) learning

Q-LEARNING

Q Factor

$$Q(s, a) = r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s})$$

$$\pi^*(s) = \operatorname{argmax}_a \left[r(s, a) + E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}) \right]$$

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

- Q is the evaluation function the agent learns

- From Q we can get V
- From V we can get Q

$$V(s) = \max_{a'} Q(s, a')$$

Q-learning

$$Q(s_t, a_t) = r(s_t, a_t) + \gamma \max_{a'} Q(s_{t+1}, a')$$

- \hat{Q} denotes learner's current approximation to Q
- At state s and action a


$$\hat{Q}(s, a) \leftarrow r(s, a) + \gamma \max_{a'} \hat{Q}(s', a')$$

- s' the state resulting from applying action a in state s
 $s' \sim p(s'|s, a)$
 - Take a single sample
- Repeat


Q-learning for Deterministic Model

- For each s, a initialize table entry $\hat{Q}(s, a) \leftarrow 0$
- Observe current state s
- Loop forever
 - Select an action a and execute it
 - Receive immediate reward r
 - Observe the new state s'
 - Update the table entry for $\hat{Q}(s, a)$ as follows
$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$
 - $s \leftarrow s'$

Q-learning



Updating \hat{Q}



$$\begin{aligned}\hat{Q}(s_1, a_{right}) &\leftarrow r + \gamma \max_{a'} \hat{Q}(s_2, a') \\ &\leftarrow 0 + 0.9 \max\{63, 81, 100\} \\ &\leftarrow 90\end{aligned}$$

- Reward non-negative

$$\begin{aligned}(\forall s, a, n) \quad \hat{Q}_{n+1}(s, a) &\geq \hat{Q}_n(s, a) \\ (\forall s, a, n) \quad 0 &\leq \hat{Q}_n(s, a) \leq Q(s, a)\end{aligned}$$

- New Q better
- Can be shown that \hat{Q} converges to Q

Nondeterministic Case

- Alter training rule to

$$\hat{Q}_n(s, a) \leftarrow (1 - \alpha_n)\hat{Q}_{n-1}(s, a) + \alpha_n \left[r + \max_{a'} \hat{Q}_{n-1}(s', a') \right]$$

- $s' \sim p(s' | s, a)$


- Learning rule

- $\alpha_n = \frac{1}{1 + \text{visits}_n(s, a)}$


- Convergence of \hat{Q} to Q still holds under mild assumptions

$$\hat{Q}_n(s, a) \leftarrow \hat{Q}_{n-1}(s, a) + \alpha_n \left[r + \max_{a'} \hat{Q}_{n-1}(s', a') - \hat{Q}_{n-1}(s, a) \right]$$


Q-learning



$a \uparrow$



$$r + \gamma \max_{a'} \hat{Q}(s', a')$$



Algorithm

- Loop
 - Sample state s
 - $a = \operatorname{argmax}_{\bar{a}} \hat{Q}(s, \bar{a})$
 - Sample $s' \sim p(s'|s, a)$
 - $\hat{Q}(s, a) \leftarrow \hat{Q}(s, a) + \alpha_n \left[r + \max_{a'} \hat{Q}(s', a') - \hat{Q}(s, a) \right]$
- Can sample more than one s' and then average
- Note
$$\sum_i \max_{a'} \hat{Q}(s'_i, a') \neq \max_{a'} \sum_i \hat{Q}(s'_i, a')$$
- $\sum_i \max_{a'} \hat{Q}(s'_i, a')$ should be used

Generalization: True Q-learning

- Loop
 - Sample state s
 - Select any action a
 - Sample $s' \sim p(s'|s, a)$
 - $\hat{Q}(s, a) \leftarrow \hat{Q}(s, a) + \alpha_n \left[r + \max_{a'} \hat{Q}(s', a') - \hat{Q}(s, a) \right]$
- Two policies
 - The selection one behind step 2
 - The 'optimal' policy behind \hat{Q}

SARSA


- Estimate value for the current action executed
- Loop
 - Execute action a at state s
 - Get reward r
 - Sample $s' \sim p(s'|s, a)$
 - $a' = \max_a \hat{Q}(s', a)$
 - $\hat{Q}(s, a) \leftarrow \hat{Q}(s, a) + \alpha_n [r + \hat{Q}(s', a') - \hat{Q}(s, a)]$
 - $s = s', a = a'$
- Just one single policy
 - Agent follows 'optimal' policy

Off-policy vs On-policy


- On-policy algorithm
 - Learn the policy being executed by the agent
- Off-policy algorithm
 - Evaluate a policy from samples generated by a different policy
 - Target policy
 - Learn policy independent of policy taken by agent
 - Behavior policy
- Q-learning is off-policy
- SARSA is on-policy

SARSA vs Q-learning

- SARSA useful when you want to optimize the value of an agent that is exploring
 - Learn near optimal while exploring
 - More conservative – slower convergence
- Do offline learning, then use the policy in an agent that does not explore
 - Q-learning more appropriate
 - Directly learns optimal policy
- Practical experiments
 - Q-values are lower in SARSA than in Q-learning
 - Does not mean that SARSA superior
 - Q-learning higher per sample variance



<http://tinyurl.com/y984s42k>



<https://sridhartee.blogspot.com/2016/09/qlearning-and-sarsa-cliff-task.html>

Value Iteration for Q-factor

- Deterministic case
- Use Bellman equation as an iterative update
$$Q_{i+1}(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q_i(s', a') | s, a \right]$$
 - Compute this for each state-action pair
- Q_i converge to Q^* as $i \rightarrow \infty$
- In practice cannot loop through all state-action pairs
 - Functional approximations must be considered

TD-LEARNING

Monte-Carlo Policy Evaluation


- Learn directly from episodes of experience
- Model-free
 - No knowledge of functions for transitions/rewards
- Learn from complete episodes
- Idea
 - Value = mean return
- Caveat
 - Applicable only to finite episodes
 - Cannot handle infinite time horizon settings

Basics

- Goal
 - Learn V^π from episodes of experience under policy π
 - $s_0, a_0, r_0, \dots, s_T \sim \pi$
- Total discounted reward
 - $G_t = r_t + \gamma r_{t+1} + \dots + \gamma^{T-1} r_T$
- Value of policy
 - $V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s]$
- Monte-Carlo policy evaluation uses empirical mean return instead of expected return


First-Visit Monte-Carlo Policy Evaluation

- To evaluate state s
 - Loop
 - Generate an episode based on π
 - The first time-step t that state s is visited in the episode
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $S(s) \leftarrow S(s) + G_t$
 - Value is estimated by mean return $V(s) = S(s)/N(s)$
- Law of large numbers
 - $V(s) \rightarrow V^\pi(s)$ as $N(s) \rightarrow \infty$



Every-Visit Monte-Carlo Policy Evaluation

- To evaluate state s
 - Loop
 - Generate an episode based on π
 - Every time-step t that state s is visited in the episode
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return $V(s) = S(s)/N(s)$
- Law of large numbers
 - $V(s) \rightarrow V^\pi(s)$ as $N(s) \rightarrow \infty$



Incremental Mean

- Mean μ_1, μ_2, \dots of sequence x_1, x_2, \dots

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

Incremental Monte-Carlo Updates

- Update $V(s)$ incrementally after episode $s_0, a_0, r_0, \dots, s_T$
- For each state s_t with return G_t

$$N(s_t) \leftarrow N(s_t) + 1$$

$$V(s_t) \leftarrow V(s_t) + \frac{1}{N(s_t)} (G_t - V(s_t))$$

- Non-stationary problems
 - Forget old episodes

$$V(s_t) \leftarrow V(s_t) + \alpha(G_t - V(s_t))$$

Temporal-Difference Learning

- Learn directly from episodes of experience
- Model-free
 - No knowledge of functions for transitions/rewards
- Learns from incomplete episodes
- Updates a guess towards a guess

MC and TD

- Goal
 - Learn V^π online from experience under policy π
- Incremental every-visit Monte-Carlo
 - Update value $V(S_t)$ toward actual return G_t
 - $V(s_t) \leftarrow V(s_t) + \alpha(G_t - V(s_t))$
- Simplest temporal-difference learning algorithm TD(0)
 - Update value $V(s_t)$ toward estimated return $r_{t+1} + \gamma V(s_{t+1})$
 - $V(s_t) \leftarrow V(s_t) + \alpha(r_t + \gamma V(s_{t+1}) - V(s_t))$
 - TD target: $r_t + \gamma V(s_{t+1})$
 - TD error: $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

TD Algorithm

- Loop
 - Generate an episode based on $\pi: s_0, a_0, r_0, \dots, s_T$
 - For $t = 0$ to T
 - $V(s_t) \leftarrow V(s_t) + \alpha(r_t + \gamma V(s_{t+1}) - V(s_t))$
- Recall from value function to policy
 - $\pi(s) = \operatorname{argmax}_a r(s, a) + E_{s' \sim p(s'|s,a)} V(s')$
 - This holds for MC and TD

MC vs. TD

- TD can learn before knowing the final outcome
 - TD can learn online after every step
 - MC must wait until end of episode before return is known
- TD can learn without the final outcome
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments

Bias/Variance Trade-Off

- Return $G_t = r_t + \gamma r_{t+1} + \dots + \gamma^{T-1} r_T$ unbiased estimate of $V^\pi(s_t)$
- True TD target $r_t + \gamma V^\pi(s_{t+1})$ unbiased estimate of $V^\pi(s_t)$
- TD target $r_t + \gamma V(s_{t+1})$ biased estimate of $V^\pi(s_t)$
- TD target much lower variance than the return
 - Return depends on many random actions, transitions, rewards
 - TD target depends on one random action, transition, reward

Advantages and Disadvantages of MC vs. TD

- MC has high variance, zero bias
 - Good convergence properties
 - Not very sensitive to initial value
 - Very simple to understand and use
- TD has low variance, some bias
 - Usually more efficient than MC
 - TD converges to $V^\pi(s)$
 - More sensitive to initial value

Batch MC and TD

- MC and TD converge
 - $V(s) \rightarrow V^\pi(s)$ as experience $\rightarrow \infty$
- Batch solution for finite experience?

$$s_0^1, a_0^1, r_0^1, \dots, s_{T_1}^1$$

\vdots

$$s_0^K, a_0^K, r_0^K, \dots, s_{T_K}^K$$

- Repeatedly sample episode $k \in [1, K]$
- Apply MC or TD to episode k

Certainty Equivalence

- MC converges to solution with minimum mean-squared error
 - Best fit to the observed returns
 - $\sum_{k=1}^K \sum_{t=0}^{T_k} (G_t^k - V(s_t^k))^2$
- TD converges to solution of max likelihood Markov model
 - Solution to the MDP $\langle \mathcal{S}, \mathcal{A}, \hat{\mathcal{P}}, \hat{\mathcal{R}}, \gamma \rangle$ that best fits the data
 - $\hat{\mathcal{P}}_{s,s'}^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=0}^{T_k} \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$
 - $\hat{\mathcal{R}}_s^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=0}^{T_k} \mathbf{1}(s_t^k, a_t^k = s, a) r_t^k$


MC vs. TD

- TD exploits Markov property
 - Usually more efficient in Markov environments
- MC does not exploit Markov property
 - Usually more efficient in non-Markov environments

$$\begin{aligned} P\{X_{t+1} = j \mid X_0 = k_0, X_1 = k_1, \dots, X_{t-1} = k_{t-1}, X_t = i\} \\ = P\{X_{t+1} = j \mid X_t = i\} \end{aligned}$$

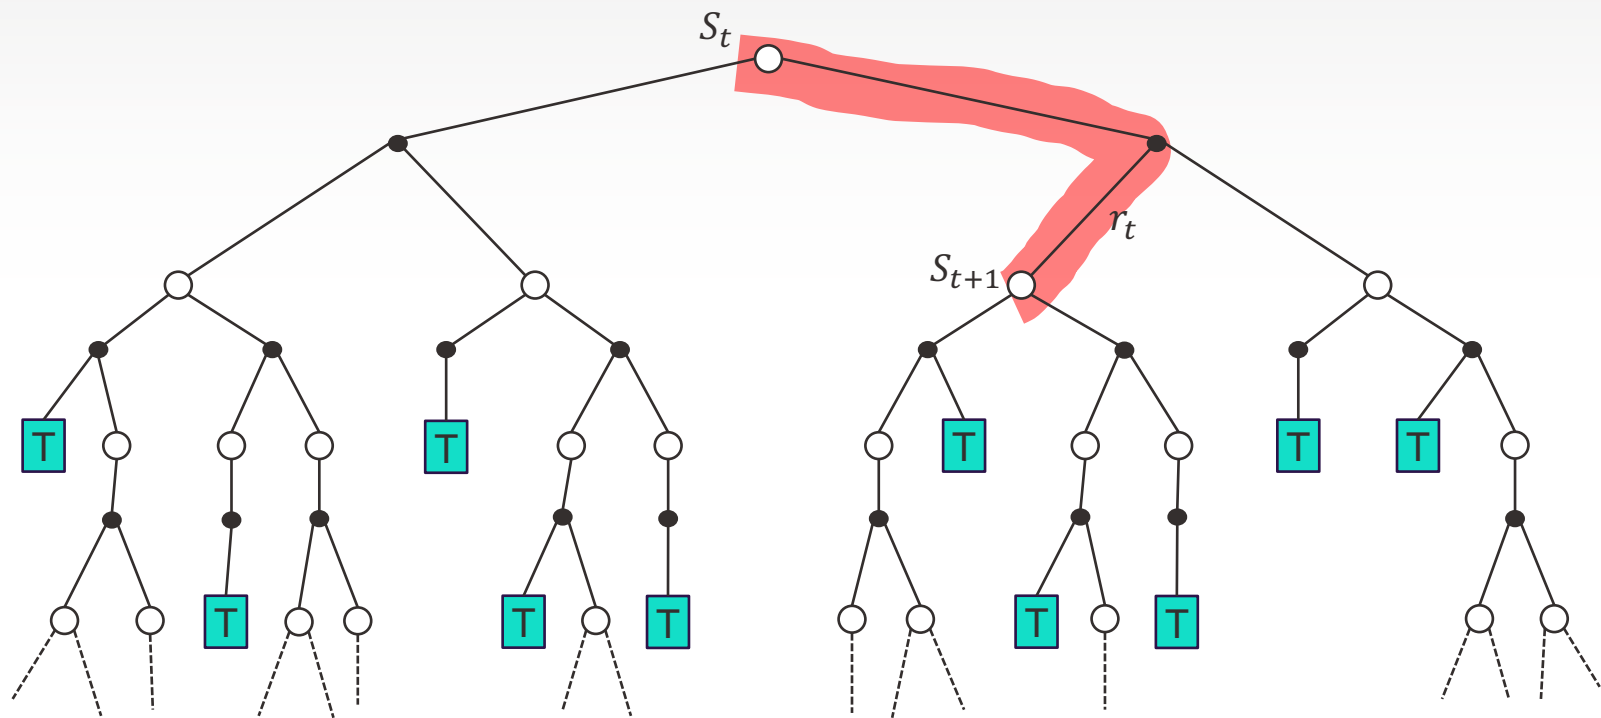
Monte-Carlo Backup

$$V(s_t) \leftarrow V(s_t) + \alpha(G_t - V(s_t))$$




Temporal-Difference Backup

$$V(s_t) \leftarrow V(s_t) + \alpha(r_t + \gamma V(s_{t+1}) - V(s_t))$$



Dynamic Programming Backup

$$V(s_t) \leftarrow \mathbb{E}_{\pi}[r_t + \gamma V(s_{t+1})]$$



n -Step Prediction

- TD can look n steps into the future

TD (1-step)



2-step



3-step



...

n -step



...

Monte-Carlo



n -Step Return

- Consider n -step returns for $n = 1, 2, \dots$

$$n = 1 \quad (TD) \quad G_t^{(1)} = r_t + \gamma V(s_{t+1})$$

$$n = 2 \quad G_t^{(2)} = r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2})$$

$$\vdots$$
$$\vdots$$

$$n = \infty \quad (MC) \quad G_t^{(\infty)} = r_t + \gamma r_{t+1} + \dots + \gamma^{T-1} r_T$$

- n -step return

$$G_t^{(n)} = r_t + \gamma r_{t+1} + \dots + \gamma^{n-1} r_{t+n-1} + \gamma^n V(s_{t+n})$$

- n -step temporal-difference learning

$$V(s) \leftarrow V(s_t) + \alpha \left(G_t^{(n)} - V(s_t) \right)$$


Averaging n -Step Returns

- Average n -step returns over different n
 - Average 2-step and 4-step returns


$$\frac{1}{2} G^{(2)} + \frac{1}{2} G^{(4)}$$

- Combines information from two different time-steps
- All of them
 - Weights should sum to 1

One backup



λ -return



- λ -return G_t^λ combines all n -step returns $G_t^{(n)}$
- Using weight $(1 - \lambda) \lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- Forward-view TD(λ)

$$V(s_t) \leftarrow V(s_t) + \alpha (G_t^\lambda - V(s_t))$$

- Original version is TD(0)

The Credit Assignment Problem

State	Reward	Action
21	0	3
40	0	2
39	0	2
38	0	2
45	0	1
44	0	1
43	100	

- Have great reward at the end
- What state-action led to high reward
- The credit assignment problem

Real-world

- Is this a purely fictitious example?
- Robots
 - Reward 0 except when hitting an object
 - Drone hitting wall otherwise reward 0
- Lose/win (zero-sum) games
 - Tic-tac-toe: reward 0 except when you have 3 in row/column/diagonal
 - Chess: reward 0 except at the end when you win or lose
 - Zero-sum: reward of player 1 = lose of player 2
- Clearly not always the case (inventory, autonomous cars)

Exploration-Exploitation

- Exploration-Exploitation tradeoff
- Have visited part of the state space and found a reward of 100
 - Is this the best we can hope for?
 - Should we keep 'pounding' the visiting states and figure out actions that lead to 100?
 - Perhaps there are other states that we have not yet visited that lead to even higher reward
- Exploitation
 - Should we stick with what we know
 - Find a good policy with respect to this knowledge
 - Risk of missing out on a better reward somewhere else
- Exploration
 - Should we look for states with more reward
 - At risk of wasting time and getting some negative reward


Eligibility Traces

- Keep in mind the assignment problem
- Possible solutions
 - Frequency heuristic
 - Assign credit to most frequent states
 - Recency heuristic
 - Assign credit to most recent states
- Eligibility traces combine both heuristics

$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(s_t = s)$$

γ = hyperparameter



accumulating eligibility trace


times of visits to a state

Backward View TD(λ)

- Keep an eligibility trace for every state s
- Update value $V(s)$ for every state s
 - In proportion to TD-error δ_t and eligibility trace $E_t(s)$

$$\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



TD(λ) and TD(0)

- When $\lambda = 0$, only current state is updated

$$E_t(s) = \mathbf{1}(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(s_t) \leftarrow V(s_t) + \alpha \delta_t$$

- Parametric value function approximations
 - Eligibility trace replaced with gradient

Algorithm

- Loop
 - For each episode $\tau = (s_0, a_0, r_0, \dots, s_T)$
 - $E = 0$
 - For $t = 0$ to T
 - For each state s
 - $E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(s_t = s)$
 - $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$
 - $V(s_t) \leftarrow V(s_t) + \alpha \delta_t E_t(s_t)$
- Alternative
 - $\delta_t = G_t^\lambda - V(s_t)$
- Note that it does not deal with policies
 - We know how to get the final policy
- Parametric approximation to E
- Updated through gradients

Connection

- The sum of offline updates is identical for forward-view and backward-view TD(λ)

$$\sum_{t=0}^T \alpha \delta_t E_t(s) = \sum_{t=0}^T \alpha \left(G_t^\lambda - V(S_t) \right) \mathbf{1}(S_t = s)$$

- Not obvious
 - Can be seen with elementary mathematics

MC and TD(1)

- Consider an episode where s is visited once at time-step k in entire episode
- TD(1) eligibility trace discounts time since visit

$$E_t(s) = \gamma E_{t-1}(s) + 1(s_t = s)$$
$$= \begin{cases} 0 & \text{if } t < k \\ \gamma^{t-k} & \text{if } t \geq k \end{cases}$$

- TD(1) updates accumulate error online

$$\sum_{t=0}^{T-1} \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^{T-1} \gamma^{t-k} \delta_t = \alpha (G_k - V(S_k))$$

- By end of episode it accumulates total error

$$\delta_k + \gamma \delta_{k+1} + \gamma^2 \delta_{k+2} + \dots + \gamma^{T-1-k} \delta_{T-1}$$

TD(1)

- Assume $\lambda = 1$
 - Sum of TD errors telescopes into MC error

$$\begin{aligned} & \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \dots + \gamma^{T-1-t}\delta_{T-1} \\ &= r_{t+1} + \gamma V(s_{t+1}) - V(s_t) \\ &+ \gamma r_{t+2} + \gamma^2 V(s_{t+2}) - \gamma V(s_{t+1}) \\ &+ \gamma^2 r_{t+3} + \gamma^3 V(s_{t+3}) - \gamma^2 V(s_{t+2}) \\ &\quad \vdots \\ &+ \gamma^{T-1-t} r_T + \gamma^{T-t} V(s_T) - \gamma^{T-1-t} V(s_{T-1}) \\ &= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} \dots + \gamma^{T-1-t} r_T - V(s_t) \\ &= G_t - V(s_t) \end{aligned}$$

- This holds also if t replaced by k
 - Shows equation on previous slide

TD(1)

- TD(1) roughly equivalent to every-visit Monte-Carlo
- Error accumulated online
 - Step-by-step
- If value function only updated offline at end of episode
 - Total update is exactly the same as MC

TD(λ)

- General λ

- $$\begin{aligned} G_t^\lambda - V(s_t) &= -V(s_t) + (1 - \lambda)\lambda^0(r_t + \gamma V(s_{t+1})) \\ &\quad + (1 - \lambda)\lambda^1(r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2})) \\ &\quad + (1 - \lambda)\lambda^2(r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 V(s_{t+3})) \\ &\quad + \dots \\ &= -V(s_t) + (\gamma\lambda)^0(r_t + \gamma V(s_{t+1}) - \gamma\lambda V(s_{t+1})) \\ &\quad + (\gamma\lambda)^1(r_{t+1} + \gamma V(s_{t+2}) - \gamma\lambda V(s_{t+2})) \\ &\quad + (\gamma\lambda)^2(r_{t+2} + \gamma V(s_{t+3}) - \gamma\lambda V(s_{t+3})) \\ &\quad + \dots \\ &= (\gamma\lambda)^0(r_t + \gamma V(s_{t+1}) - V(s_t)) \\ &\quad + (\gamma\lambda)^1(r_{t+1} + \gamma V(s_{t+2}) - V(s_{t+1})) \\ &\quad + (\gamma\lambda)^2(r_{t+2} + \gamma V(s_{t+3}) - V(s_{t+2})) \\ &\quad + \dots \\ &= \delta_t + \gamma\lambda\delta_{t+1} + (\gamma\lambda)^2\delta_{t+2} + \dots \end{aligned}$$

Summary

- TD(0) is the original version
 - With eligibility traces
 - One step look ahead
- TD(1)
 - Look ahead all the way to the end of an episode
 - But from the first occurrence of state
 - Very similar to MC
- TD(λ)
 - Errors weighted by $\gamma\lambda$ all the way through the end of an episode